

SALEM NASSAR

Austin, TX • 512-202-1126 • Salem.Nassar1@outlook.com
linkedin.com/in/salemnassar/ • github.com/Salemmander • salemnassar.com

Profile

Software engineer with 2+ years at Verizon building Python automation and OpenShift infrastructure that runs in production. Outside of work, I build autonomous systems in PyTorch and C++, including a self-driving RC car that runs C++ inference on a Raspberry Pi. Currently pursuing an M.S. in Computer Science with a focus in Machine Learning at Georgia Tech.

Experience

Software Engineer

June 2024 – Present

Verizon, Southlake, TX

- Built 20+ Python automations that generate and send router configs for sites and access leaves nationwide across 8+ vendor platforms (Cisco, Nokia, Ericsson, Samsung, Ribbon)
- Reduced deployment time by 90% with GitLab CI/CD multi-stage Docker builds and OpenTofu-automated multi-cluster OpenShift deployments
- Solely deployed Nautobot and HashiCorp Vault to 5 sites with OpenTofu, automating multi-site OpenShift rollouts for both

Projects

Autonomous RC Car

- Built an end-to-end self-driving RC car on Raspberry Pi 5, implementing NVIDIA's **PilotNet CNN architecture** ([Bojarski et al., 2016](#)) in PyTorch to predict steering and throttle from raw camera images via behavioral cloning
- Trained the 5-layer CNN on **10K+ frames** with early stopping, enabling the car to navigate the track **100% autonomously** in both directions
- Built a multithreaded **C++** runtime with a dedicated OpenCV camera thread and **ONNX Runtime** inference in **3.6 ms** on Pi 5 CPU

Reinforcement Learning Racing Agent (AWS DeepRacer)

- Implemented **PPO** in PyTorch for continuous steering and throttle control, trained across 4 parallel simulators
- Designed a **sensor-fusion** policy that combines stereo cameras (2D CNN) and a 64-ray **LiDAR** scan (1D CNN) into shared actor-critic features, with no access to ground-truth position or heading
- Completed **15/15** official Time-Trial evaluation episodes across 3 tracks

Technical Skills

- **Machine Learning and Computer Vision:** PyTorch, ONNX, NumPy, OpenCV, CNNs, Behavioral Cloning, Reinforcement Learning (PPO), Scikit-learn
- **Programming Languages:** Python, C++, C, Bash
- **DevOps and Automation Tools:** Kubernetes, OpenShift, Docker, Helm, GitLab CI/CD, OpenTofu/Terraform, Linux, Git, Vault
- **Databases:** PostgreSQL, SQL, Redis, MongoDB

Education

Master of Science in Computer Science, Specialization in Machine Learning

Aug 2025 – May 2028

Georgia Institute of Technology (Online)

- GPA: 4.0
- Completed: Artificial Intelligence, Computer Vision, Machine Learning, AI Techniques for Robotics (SLAM, localization, motion planning), Reinforcement Learning
- Current: Deep Learning

Bachelor of Science in Computer Science

May 2024

The University of Texas at Dallas

- Relevant Coursework: Data Structures and Algorithms, Machine Learning, Artificial Intelligence